

learning curves: SGD, one hidden layer NN
overparametrized
input dim: 50, n_samples size: 20
hidden dim: 80
output dim: 1

